

0.1 Project #3:

- Student name: Milena Afeworki
- Student pace: full time
- Scheduled project review date/time: 07/09/2021 @ 10:15 PT
- Instructor name: Abhineet Kulkarni
- Blog post URL:

In []: ▶

1 Tanzanian Water Well Functionality Classification

Tanzania, as a developing country, struggles with providing clean water to its population of over 57,000,000. There are many waterpoints already established in the country, but some are in need of repair while others have failed altogether. Using data from Taarifa and the Tanzanian Ministry of Water, we need to predict which pumps are functional, which need some repairs, and which don't work at all. A smart understanding of which waterpoints will fail can improve maintenance operations and ensure that clean, potable water is available to communities across Tanzania.

1.1 The Business Problem

This project aims to use the data to anticipate when a well needs repair or maintenance, ideally before it breaks and disrupts the local water supply. Failing to identify nonfunctional water supply lines could lead villagers to suffer in many ways, including traveling long distances to other water sources resulting in increased time and effort to fetch water, and being exposed to different health-related issues that come with poor water quality. Accordingly, though it is important to build a model that will accurately classify the wells, it is crucial that our model's tolerance to errors of misclassifying wells, especially the 'nonfunctional' and 'needs repair' groups, is as low as possible.

1.2 Data Understanding

- **amount_tsh** - Total static head (amount water available to waterpoint)
- **date_recorded** - The date the row was entered



- **funder** - Who funded the well
- **gps_height** - Altitude of the well
- **installer** - Organization that installed the well
- **longitude** - GPS coordinate
- **latitude** - GPS coordinate
- **wpt_name** - Name of the waterpoint if there is one
- **num_private** -
- **basin** - Geographic water basin
- **subvillage** - Geographic location
- **region** - Geographic location
- **region_code** - Geographic location (coded)
- **district_code** - Geographic location (coded)
- **lga** - Geographic location
- **ward** - Geographic location
- **population** - Population around the well
- **public_meeting** - True/False
- **recorded_by** - Group entering this row of data
- **scheme_management** - Who operates the waterpoint
- **scheme_name** - Who operates the waterpoint
- **permit** - If the waterpoint is permitted
- **construction_year** - Year the waterpoint was constructed
- **extraction_type** - The kind of extraction the waterpoint uses
- **extraction_type_group** - The kind of extraction the waterpoint uses
- **extraction_type_class** - The kind of extraction the waterpoint uses
- **management** - How the waterpoint is managed
- **management_group** - How the waterpoint is managed
- **payment** - What the water costs
- **payment_type** - What the water costs
- **water_quality** - The quality of the water
- **quality_group** - The quality of the water
- **quantity** - The quantity of water
- **quantity_group** - The quantity of water
- **source** - The source of the water
- **source_type** - The source of the water
- **source_class** - The source of the water
- **waterpoint_type** - The kind of waterpoint
- **waterpoint_type_group** - The kind of waterpoint

```
In [1]: ▶ # import all the necessary Libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import train_test_split, cross_val_score

from sklearn.metrics import confusion_matrix, classification_report
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
# sns.set_style('whitegrid')

from sklearn.metrics import plot_confusion_matrix

from xgboost import XGBClassifier
```

executed in 1.32s, finished 08:10:04 2021-07-11

1.3 Obtain data

Let's first import the data and take a look at the info to see if we need to do some data cleaning.

```
In [2]: # Load data set
df = pd.read_csv('training_set_values.csv')
df.head()
```

executed in 335ms, finished 08:10:04 2021-07-11

Out[2]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name	num_private	...	payment_ty
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	none	0	...	annu
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati	0	...	never p
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi	0	...	per buc
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	...	never p
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Shuleni	0	...	never p

5 rows × 40 columns



In [3]:  df.info()

executed in 78ms, finished 08:10:04 2021-07-11

```
22 permit                56344 non-null object
23 construction_year      59400 non-null int64
24 extraction_type        59400 non-null object
25 extraction_type_group   59400 non-null object
26 extraction_type_class   59400 non-null object
27 management             59400 non-null object
28 management_group        59400 non-null object
29 payment                59400 non-null object
30 payment_type            59400 non-null object
31 water_quality           59400 non-null object
32 quality_group           59400 non-null object
33 quantity               59400 non-null object

34 quantity_group         59400 non-null object
35 source                 59400 non-null object
36 source_type            59400 non-null object
37 source_class           59400 non-null object
38 waterpoint_type        59400 non-null object
39 waterpoint_type_group   59400 non-null object
dtypes: float64(3), int64(7), object(30)
```

```
In [4]: # check for NaNs  
df.isna().sum()
```

executed in 79ms, finished 08:10:04 2021-07-11

```
Out[4]: id                0  
amount_tsh              0  
date_recorded           0  
funder                 3635  
gps_height              0  
installer              3655  
longitude               0  
latitude                0  
wpt_name                0  
num_private             0  
basin                   0  
subvillage              371  
region                  0  
region_code             0  
district_code           0  
lga                     0  
ward                    0  
population              0  
public_meeting          3334  
recorded_by             0  
scheme_management       3877  
scheme_name             28166  
permit                  3056  
construction_year       0  
extraction_type         0  
extraction_type_group   0  
extraction_type_class   0  
management              0  
management_group        0  
payment                 0  
payment_type            0  
water_quality           0  
quality_group           0  
quantity                0  
quantity_group          0  
source                  0  
source_type             0  
source_class            0
```

```
waterpoint_type      0  
waterpoint_type_group 0  
dtype: int64
```

1.4 Scrubbing the data

1.4.1 Cleaning based on info

Key observations from here:

1. Dealing with missing values:

- funder = 3635
- installer = 3655
- subvillage = 371
- public_meeting = 3334
- scheme_managment = 3877
- scheme_name = 28166
- permit = 3056

2. Dealing with date_recorded data type.

3. Dealing with outliers.

1.4.2 Dealing with the missing values.

Funder

```
In [5]: df.funder.value_counts(normalize=True)
```

executed in 30ms, finished 08:10:04 2021-07-11

```
Out[5]: Government Of Tanzania    0.162898
Danida                          0.055841
Hesawa                          0.039487
Rwssp                           0.024639
World Bank                      0.024191
...
Kanisa La Neema                 0.000018
Tcrst                          0.000018
Raramataki                     0.000018
Tasae                          0.000018
Africaone Ltd                   0.000018
Name: funder, Length: 1897, dtype: float64
```

```
In [6]: # Replace NaNs in 'funder' by 'other'
df['funder'] = df['funder'].replace(np.nan, 'other')
```

executed in 12ms, finished 08:10:04 2021-07-11

```
In [7]: df.drop(columns=['scheme_name', 'subvillage', 'public_meeting',
                        'num_private', 'permit'], axis=1, inplace = True)
```

executed in 31ms, finished 08:10:04 2021-07-11

Scheme managment


```
In [8]: df.scheme_management.value_counts()
```

executed in 13ms, finished 08:10:04 2021-07-11

```
Out[8]: VWC          36793
        WUG           5206
        Water authority 3153
        WUA           2883
        Water Board    2748
        Parastatal     1680
        Private operator 1063
        Company        1061
        Other          766
        SWC            97
        Trust          72
        None           1
        Name: scheme_management, dtype: int64
```

```
In [9]: # Replace NaNs in 'scheme_managment' by 'other'
df['scheme_management'] = df['scheme_management'].replace(np.nan, 'other')
```

executed in 14ms, finished 08:10:04 2021-07-11

We might need this feature as it might give a better glance at which organization is responsible for the management of a water well project scheme.

Installer

```
In [10]: # Replace NaNs in 'installer' by 'other'
df['installer'] = df['installer'].replace(np.nan, 'other')
```

executed in 14ms, finished 08:10:04 2021-07-11

```
In [ ]:
```

```
In [ ]:
```

1.4.3 Dealing with date_recorded

Tanzania has a rainy/wet season from December to May and a dry season from July to October. But as seen from the data not all the wells receive their water source from rainfall so the season may not be of importance to us. Also the date recorded doesn't really signify age of the well so we are going to drop it all together.

```
In [11]: df.drop(columns=['date_recorded'], axis=1, inplace = True)
```


executed in 31ms, finished 08:10:04 2021-07-11

1.4.4 Dealing with outliers

Lets take a look at the location of those wells on the map and explore for any misplaced data or anything that looks weird.

```
In [12]: sns.scatterplot('latitude', 'longitude', data=df);
```

executed in 263ms, finished 08:10:05 2021-07-11



Looking at the scatter plot of the locational coordinates, we notice an outlier with a 0' 0' latitude and longitude which really doesn't make sense since these points are far off Tanzania. In this next step lets see how many of our data have these coordinates and drop them accordingly.

```
In [13]: ► test = df.loc[df['longitude']==0, 'latitude'].value_counts()  
test
```

executed in 14ms, finished 08:10:05 2021-07-11

```
Out[13]: -2.000000e-08    1812  
Name: latitude, dtype: int64
```

```
In [14]: ► to_drop = df.loc[(df['longitude']==0) & (df['latitude'] == -2.000000e-08)]  
to_drop.shape
```

executed in 14ms, finished 08:10:05 2021-07-11

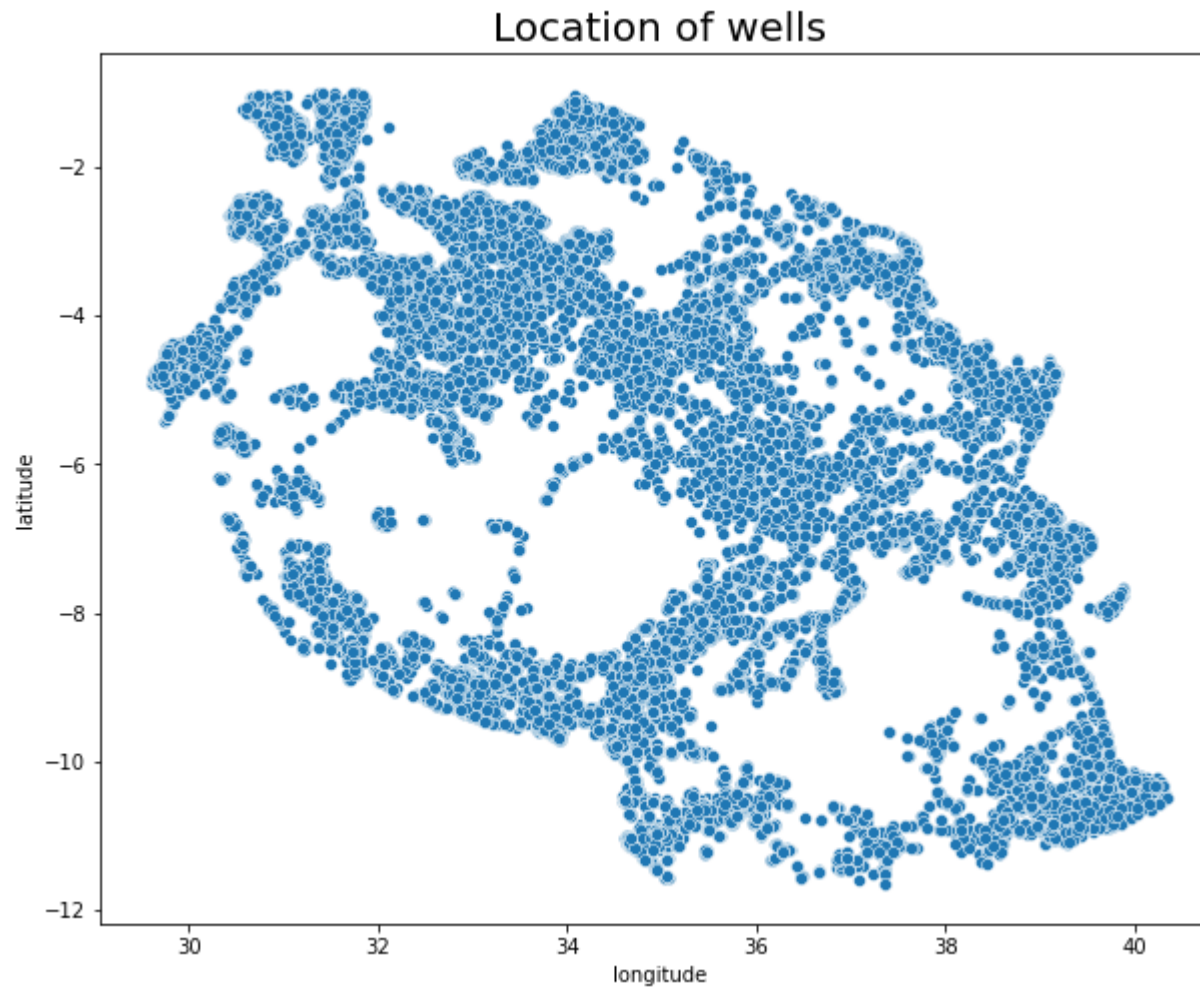
```
Out[14]: (1812, 34)
```

```
In [15]: ► df.drop(to_drop.index, axis=0, inplace=True)
```

executed in 31ms, finished 08:10:05 2021-07-11

```
In [16]: #plot Longitude and latitude of Tanzania  
fig = plt.figure(figsize=(10,8))  
sns.scatterplot('longitude', 'latitude', data=df)  
plt.title('Location of wells', fontsize=20);
```

executed in 268ms, finished 08:10:05 2021-07-11



Now this looks much better. Check one last time to see if we have any missing values.

In [17]: `df.info()`

executed in 77ms, finished 08:10:05 2021-07-11

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 57588 entries, 0 to 59399
Data columns (total 34 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    57588 non-null  int64
1   amount_tsh            57588 non-null  float64
2   funder                57588 non-null  object
3   gps_height            57588 non-null  int64
4   installer             57588 non-null  object
5   longitude             57588 non-null  float64
6   latitude              57588 non-null  float64
7   wpt_name              57588 non-null  object
8   basin                57588 non-null  object
9   region                57588 non-null  object
10  region_code           57588 non-null  int64
11  district_code         57588 non-null  int64
12  lga                   57588 non-null  object
13  ward                  57588 non-null  object
14  population            57588 non-null  int64
15  recorded_by           57588 non-null  object
16  scheme_management     57588 non-null  object
17  construction_year     57588 non-null  int64
18  extraction_type       57588 non-null  object
19  extraction_type_group  57588 non-null  object
20  extraction_type_class  57588 non-null  object
21  management            57588 non-null  object
22  management_group      57588 non-null  object
23  payment               57588 non-null  object
24  payment_type          57588 non-null  object
25  water_quality         57588 non-null  object
26  quality_group         57588 non-null  object
27  quantity              57588 non-null  object
28  quantity_group        57588 non-null  object
29  source                57588 non-null  object
30  source_type           57588 non-null  object
31  source_class          57588 non-null  object
32  waterpoint_type       57588 non-null  object
33  waterpoint_type_group 57588 non-null  object

```

```
dtypes: float64(3), int64(6), object(25)  
memory usage: 15.4+ MB
```

1.5 Explore

Now that our data is clean we will move on to the next step and merge the two tables to their corresponding ids inorder to label them according to their functionality.

1.5.1 Merging Labels to the well ids

```
In [18]: # Read csv file  
df1 = pd.read_csv('training_set_labels.csv')  
df1.head()
```

executed in 29ms, finished 08:10:05 2021-07-11

Out[18]:

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional

```
In [19]: # merge the two tables  
df2 = pd.merge(df, df1, how = 'inner', left_on = ['id'], right_on = ['id'])  
df2.shape
```

executed in 139ms, finished 08:10:05 2021-07-11

Out[19]: (57588, 35)

```
In [20]: df2.status_group.value_counts()
```

executed in 15ms, finished 08:10:05 2021-07-11

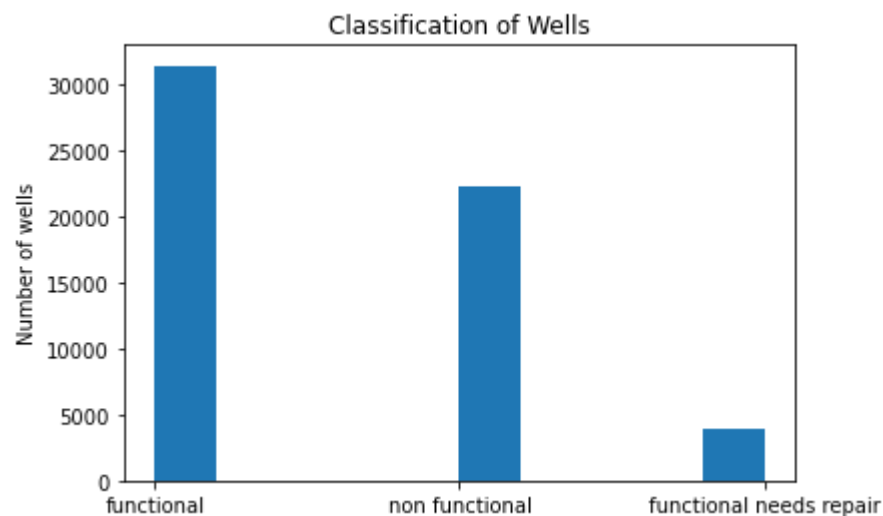
```
Out[20]: functional          31389  
non functional             22268  
functional needs repair    3931  
Name: status_group, dtype: int64
```

1.5.2 Visualization of features in relation to functionality

```
In [21]: #plot number of wells according to functionality  
plt.hist('status_group', data=df2);  
plt.title('Classification of Wells')  
plt.ylabel('Number of wells')
```

executed in 141ms, finished 08:10:06 2021-07-11

```
Out[21]: Text(0, 0.5, 'Number of wells')
```




```
In [22]: df2.status_group.value_counts(normalize=True)
```

```
executed in 15ms, finished 08:10:06 2021-07-11
```

```
Out[22]: functional          0.545061  
non functional             0.386678  
functional needs repair    0.068261  
Name: status_group, dtype: float64
```

We do see a class imbalance in the status group with 54.5% functional, 38.67% non functional, 6.82% functional needs repair.

```
In [23]: ▶ # plot wells on map with respect to water point height
plt.figure(figsize=(12,10))

plt.scatter(x='longitude', y='latitude', c='gps_height' , data=df2, s=10, cmap='twilight')
plt.colorbar().set_label('GPS Height')
plt.xlabel('Longitude', fontsize=15)
plt.ylabel('Latitude', fontsize=15)
plt.title('Location of Wells and their GPS height', fontsize=20)

plt.show()
```

executed in 798ms, finished 08:10:06 2021-07-11

Expectedly, wells with low GPS height of their water point seem to be clustered around 'non functional' or 'functional needs repair' classes. We will take a look at those in the actual map below and try to see the relationship.

```
In [24]: ▶ import cartopy.crs as ccrs
import cartopy.feature as cfeature

plt.figure(figsize=(14,10))

# Creates the map
ca_map = plt.axes(projection=ccrs.PlateCarree())


ca_map.add_feature(cfeature.LAND)
ca_map.add_feature(cfeature.OCEAN)
ca_map.add_feature(cfeature.COASTLINE)
ca_map.add_feature(cfeature.BORDERS, linestyle='-')
ca_map.add_feature(cfeature.LAKES)
ca_map.add_feature(cfeature.RIVERS, linestyle='-')
# ca_map.add_feature(cfeature.STATES.with_scale('10m'))

# Plots the data onto map
sns.scatterplot(df2['longitude'], df2['latitude'],
               s=30,
               hue=df2['status_group'],
               transform=ccrs.PlateCarree())

# Plot Labels
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)
plt.title('Functionality Status of Wells', fontsize=25)
plt.legend()
plt.show()
```

executed in 7.44s, finished 08:10:14 2021-07-11

Functionality Status of Wells



```
In [25]: ▶ import cartopy.crs as ccrs
import cartopy.feature as cfeature

plt.figure(figsize=(14,10))

# Creates the map
ca_map = plt.axes(projection=ccrs.PlateCarree())

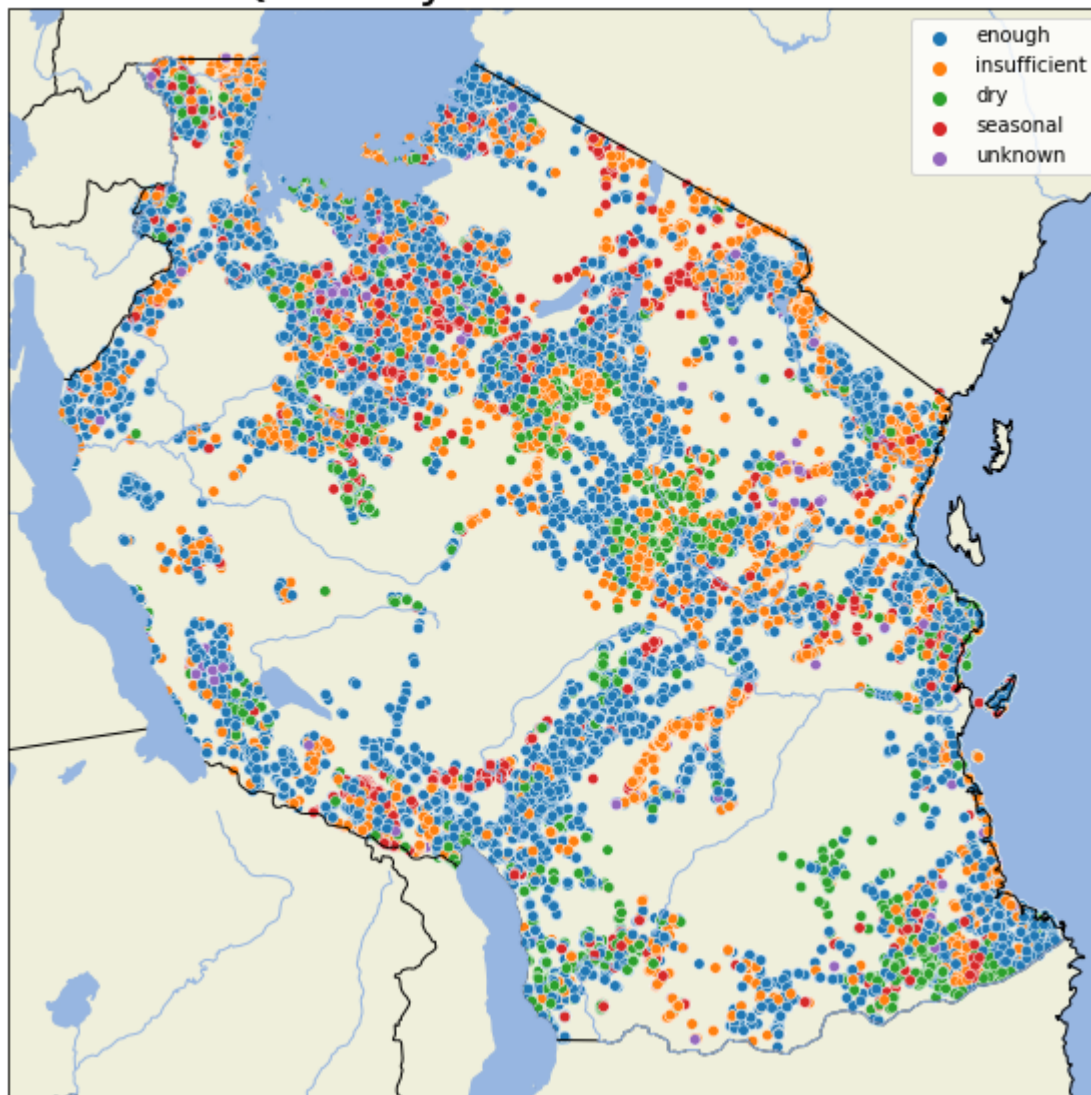
ca_map.add_feature(cfeature.LAND)
ca_map.add_feature(cfeature.OCEAN)
ca_map.add_feature(cfeature.COASTLINE)
ca_map.add_feature(cfeature.BORDERS, linestyle='-')
ca_map.add_feature(cfeature.LAKES)
ca_map.add_feature(cfeature.RIVERS)
# ca_map.add_feature(cfeature.STATES.with_scale('10m'))

# Plots the data onto map
sns.scatterplot(df2['longitude'], df2['latitude'],
               s=30,
               hue=df2['quantity'],
               transform=ccrs.PlateCarree())

# Plot Labels
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)
plt.title('Quantity of water in Wells', fontsize=25)
plt.legend()
plt.show()
```

executed in 1.65s, finished 08:10:15 2021-07-11

Quantity of water in Wells



```
In [26]: ▶ df.amount_tsh.value_counts(normalize=True)
```

```
executed in 14ms, finished 08:10:15 2021-07-11
```

```
Out[26]: 0.0          0.691585
         500.0        0.053865
         50.0         0.042926
        1000.0        0.025839
         20.0         0.025405
         ...
        8500.0        0.000017
        6300.0        0.000017
         220.0        0.000017
       138000.0        0.000017
          12.0        0.000017
        Name: amount_tsh, Length: 98, dtype: float64
```

Majority of the wells (69%) have a total static head of 0, which could signify that they are dry. Since this data was recorded at a specific date though, this could mean that those wells are non functional or possibly seasonal and happened to be dry at the time of inspection.

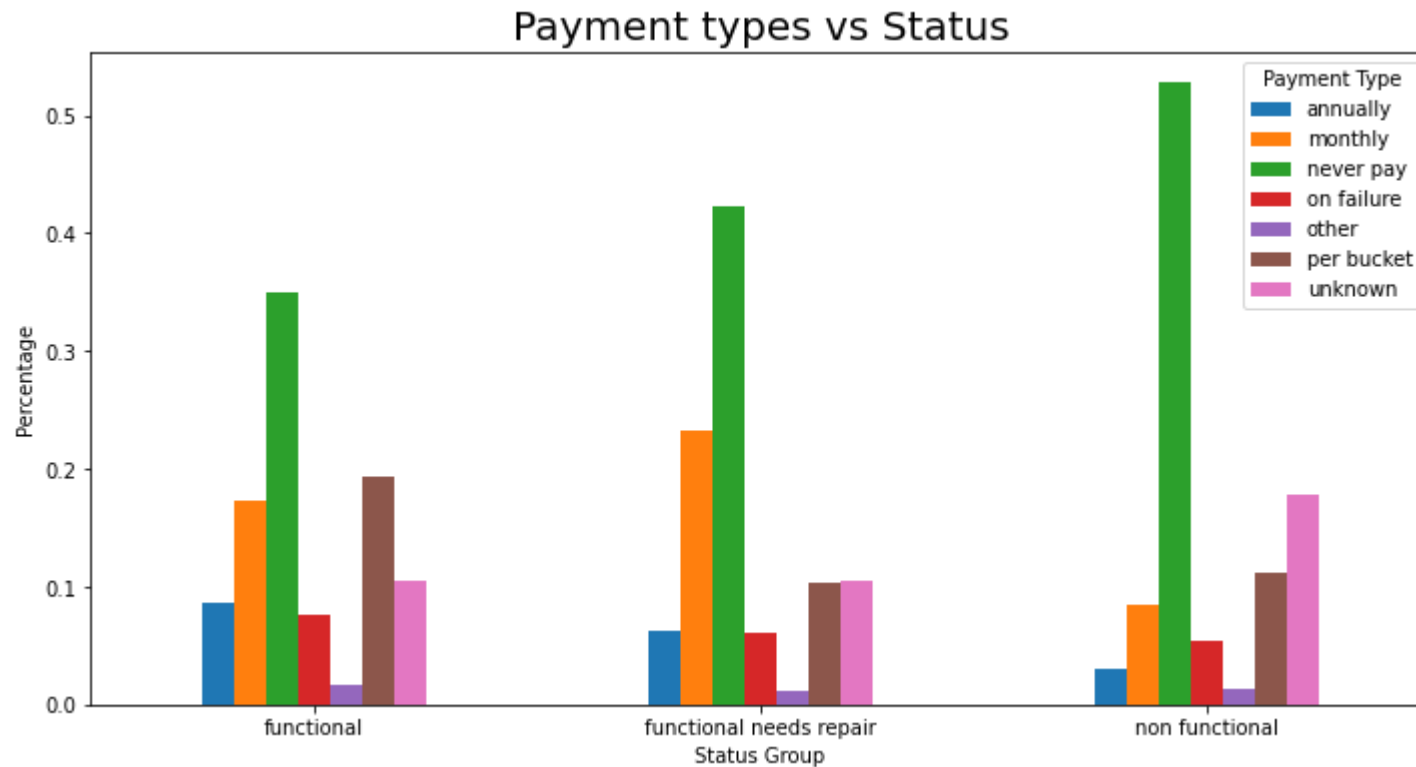
```
In [27]: ▶ payment = df2.groupby('status_group')['payment_type'].value_counts(normalize = True).unstack()
```

```
executed in 30ms, finished 08:10:15 2021-07-11
```

```
In [28]: payment.plot.bar(figsize = (12, 6))
plt.xticks(rotation = 0)
plt.title('Payment types vs Status', fontsize=20)
plt.legend(title = 'Payment Type')
plt.xlabel('Status Group')
plt.ylabel('Percentage')
```

executed in 286ms, finished 08:10:16 2021-07-11

Out[28]: Text(0, 0.5, 'Percentage')




Let's keep plotting visualizations of different features with respect to status group of wells to understand the trend of the functionality of the those wells


```
In [29]: ▶ quantity_df = df2.groupby('quantity')['status_group'].value_counts().unstack();
```

executed in 29ms, finished 08:10:16 2021-07-11

```
In [30]: ▶ quantity_df.plot.bar(figsize = (10, 6), width=0.7)
plt.title('Status vs quantity in wells', fontsize=20)
plt.xlabel('Quantity in wells', fontsize=15)
plt.ylabel('Number of wells', fontsize=15)
plt.xticks(rotation = 0);
```

executed in 283ms, finished 08:10:16 2021-07-11



In [31]: `df2.quantity.value_counts(normalize=True)`

executed in 12ms, finished 08:10:16 2021-07-11

Out[31]:

enough	0.560186
insufficient	0.252900
dry	0.104015
seasonal	0.069476
unknown	0.013423

Name: quantity, dtype: float64

In [32]: `quantity_df = df2.groupby('water_quality')['status_group'].value_counts().unstack()`

executed in 30ms, finished 08:10:16 2021-07-11

In [33]: `quantity_df.plot.bar(figsize = (12, 6), width=0.7)
plt.title('Status vs Quality of wells')
plt.xlabel('Quality of wells')
plt.ylabel('Number of wells')
plt.xticks(rotation = 0);`

executed in 253ms, finished 08:10:16 2021-07-11

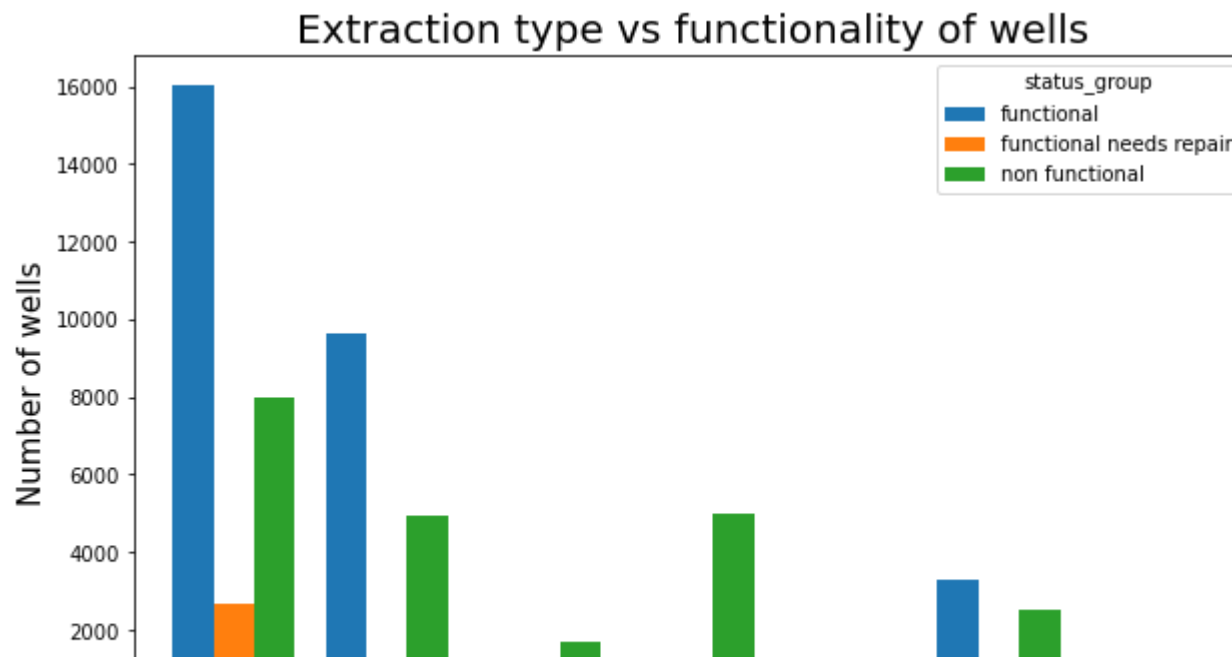
```
In [34]: quantity_df = df2.groupby('extraction_type_class')['status_group'].value_counts().unstack()
```

executed in 29ms, finished 08:10:16 2021-07-11

```
In [35]: quantity_df.plot.bar(figsize = (10, 6), width=0.8)

plt.title('Extraction type vs functionality of wells', fontsize=(20))
plt.xlabel('Extraction Type', fontsize=15)
plt.ylabel('Number of wells', fontsize=15)
plt.xticks(rotation = 0);
```

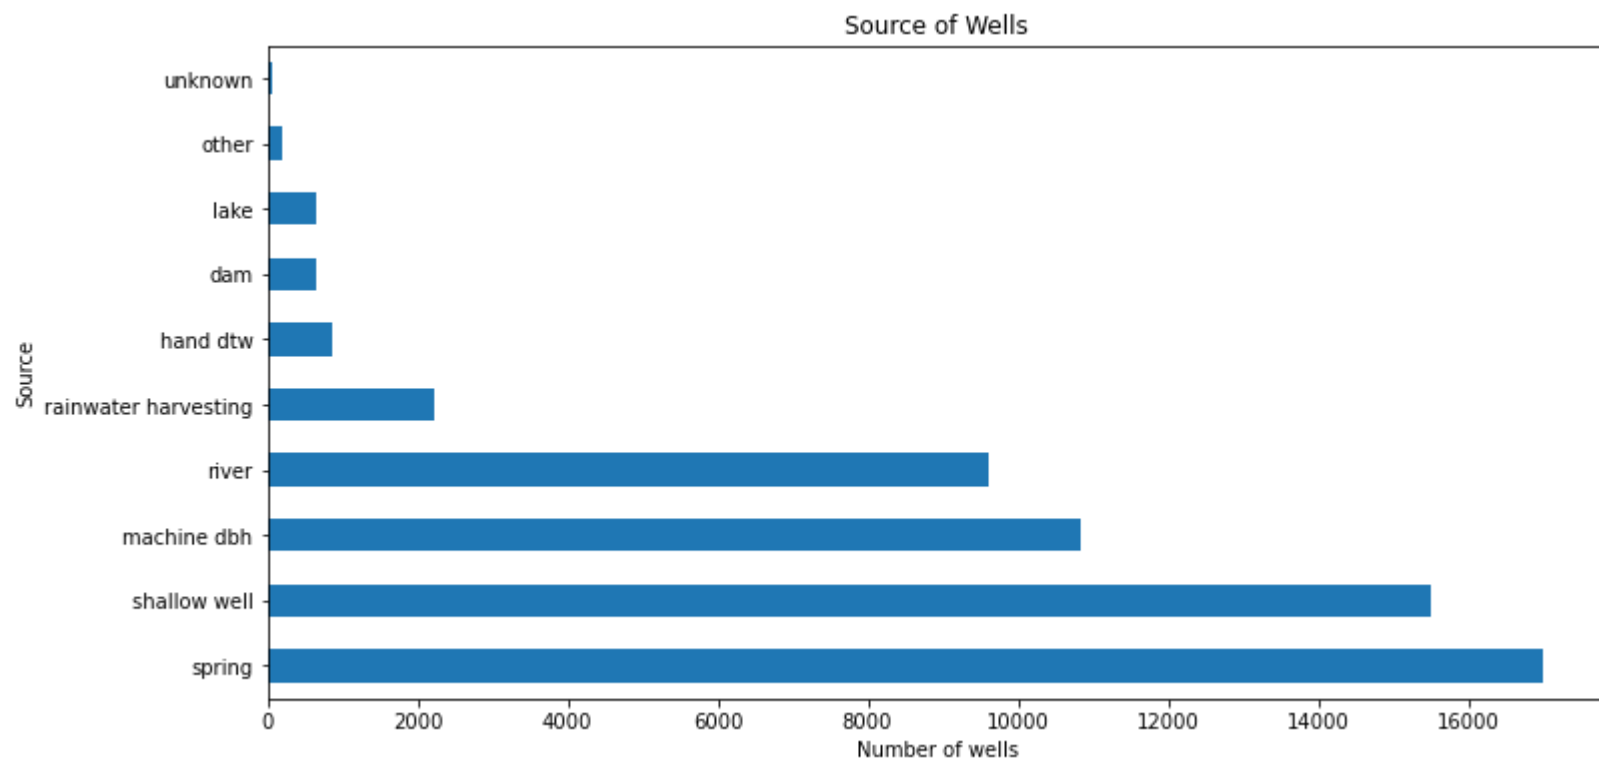
executed in 252ms, finished 08:10:17 2021-07-11



```
In [36]: df2.source.value_counts().plot.barh(figsize = (12, 6))  
plt.title('Source of Wells')  
plt.ylabel('Source')  
plt.xlabel('Number of wells')
```

executed in 218ms, finished 08:10:17 2021-07-11

Out[36]: Text(0.5, 0, 'Number of wells')




```
In [37]: ▶ plt.figure(figsize=(14,10))
sns.scatterplot(x='longitude', y='latitude', hue='water_quality' , data=df2)

plt.xlabel('Longitude', fontsize=15)
plt.ylabel('Latitude', fontsize=15)
plt.title('Location of Wells and their GPS height', fontsize=20)

plt.show()
```

executed in 1.44s, finished 08:10:18 2021-07-11



1.5.3 Construction year

Construction year is an important feature for our modeling but we do have a large range of years which will be a huge number of columns when creating dummies. We will categorize the years in such a way that they will be parts of a 10 year time period to reduce the classes of this feature. Another point worth mentioning is that more than 30% of our data has no record of the construction year. Since this is an

important feature and losing that much amount of data is not worth the risk we will take the option of replacing those values with random choice from the rest of the data.

```
In [38]: df2.construction_year.value_counts(normalize=True)
```

```
executed in 14ms, finished 08:10:18 2021-07-11
```

```
Out[38]: 0      0.328141
2010    0.045930
2008    0.045374
2009    0.043985
2000    0.036310
2007    0.027558
2006    0.025544
2003    0.022331
2011    0.021810
2004    0.019501
2012    0.018823
2002    0.018667
1978    0.018007
1995    0.017608
2005    0.017556
1999    0.017000
1998    0.016774
1990    0.016566
1985    0.016410
1980    0.014083
1996    0.014083
1984    0.013527
1982    0.012919
1994    0.012815
1972    0.012294
1974    0.011739
1997    0.011183
1992    0.011113
1993    0.010558
2001    0.009377
1988    0.009047
1983    0.008474
1975    0.007588
1986    0.007536
1976    0.007189
1970    0.007137
1991    0.005626
1989    0.005487
1987    0.005244
```



```
1981    0.004133
1977    0.003508
1979    0.003334
1973    0.003195
2013    0.003056
1971    0.002518
1960    0.001771
1967    0.001528
1963    0.001476
1968    0.001337
1969    0.001025
1964    0.000695
1962    0.000521
1961    0.000365
1965    0.000330
1966    0.000295
```

Name: construction_year, dtype: float64

```
In [39]: df2['construction_year'] = df2['construction_year'].replace(0, np.nan)
```

executed in 14ms, finished 08:10:18 2021-07-11

```
In [40]: #replace missing values with random choice
s = df2.construction_year.value_counts(normalize=True)
df2['const_year'] = df2['construction_year']
df2.loc[df2.construction_year.isna(), 'const_year'] = np.random.choice(s.index, p=s.values, size=df2.constru
```

executed in 46ms, finished 08:10:18 2021-07-11

```
In [41]: df2['const_year'].value_counts(normalize = True)
```

executed in 19ms, finished 08:10:18 2021-07-11

```
Out[41]: 2010.0    0.069077
2008.0    0.067549
2009.0    0.064823
2000.0    0.054716
2007.0    0.041311
2006.0    0.038428
2003.0    0.032906
2011.0    0.032472
2004.0    0.028756
2002.0    0.028079
2012.0    0.027870
1978.0    0.027106
2005.0    0.026117
1995.0    0.025682
1998.0    0.025526
1999.0    0.025023
1985.0    0.024693
1990.0    0.024363
1980.0    0.020768
1996.0    0.020661
```

```
In [42]: df2.drop('construction_year', axis=1, inplace=True)
```

executed in 30ms, finished 08:10:18 2021-07-11

Now that we will have replaced the missing values, next step would be to create bins for the years. The typical life expectancy of a water well is supposedly 65-100 years, and the life expectancy of a water well pump is 10-15 years. We will bin the construction year column in such a way and then take a look at the visualizations.

```
In [43]: norm = df2.const_year.value_counts()
```

executed in 13ms, finished 08:10:18 2021-07-11

```
In [44]: norm = df2.groupby('const_year')['status_group'].value_counts().unstack()
```


executed in 29ms, finished 08:10:18 2021-07-11

In [45]:

```
norm.plot.barh(figsize = (12, 10))  
plt.title('Pump by time period')  
plt.xlabel('pump status')  
plt.ylabel('number of pumps')  
plt.xticks(rotation = 0)  
plt.legend(title = 'Time Periods')
```

executed in 1.50s, finished 08:10:20 2021-07-11

Out[45]: <matplotlib.legend.Legend at 0x1ef8b060fd0>



```
In [46]: bins = [1960, 1970, 1980, 1990, 2000, 2010, 2015]
periods = ['1960-1969', '1970-1979', '1980-1989', '1990-1999', '2000-2009', '2010-2015']

df2['construction_periods'] = pd.cut(df2['const_year'], bins,
                                     labels = periods, right = False)
```

executed in 13ms, finished 08:10:20 2021-07-11

```
In [47]: df2.head()
```

executed in 28ms, finished 08:10:20 2021-07-11

Out[47]:

	id	amount_tsh	funder	gps_height	installer	longitude	latitude	wpt_name	basin	region	...	quantity	quantity_
0	69572	6000.0	Roman	1390	Roman	34.938093	-9.856322	none	Lake Nyasa	Iringa	...	enough	e
1	8776	0.0	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati	Lake Victoria	Mara	...	insufficient	insuf
2	34310	25.0	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi	Pangani	Manyara	...	enough	e
3	67743	0.0	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	Ruvuma / Southern Coast	Mtwara	...	dry	
4	19728	0.0	Action In A	0	Artisan	31.130847	-1.825359	Shuleni	Lake Victoria	Kagera	...	seasonal	se

5 rows × 36 columns



In [48]: `df2.construction_periods.value_counts()`

executed in 15ms, finished 08:10:20 2021-07-11

Out[48]:

2000-2009	22849
1990-1999	11384
1980-1989	8223
2010-2015	7730
1970-1979	6603
1960-1969	799

Name: construction_periods, dtype: int64

In [49]: `time_periods = df2.groupby('status_group')['construction_periods'].value_counts().unstack()
time_periods.head()`

executed in 31ms, finished 08:10:20 2021-07-11

Out[49]:

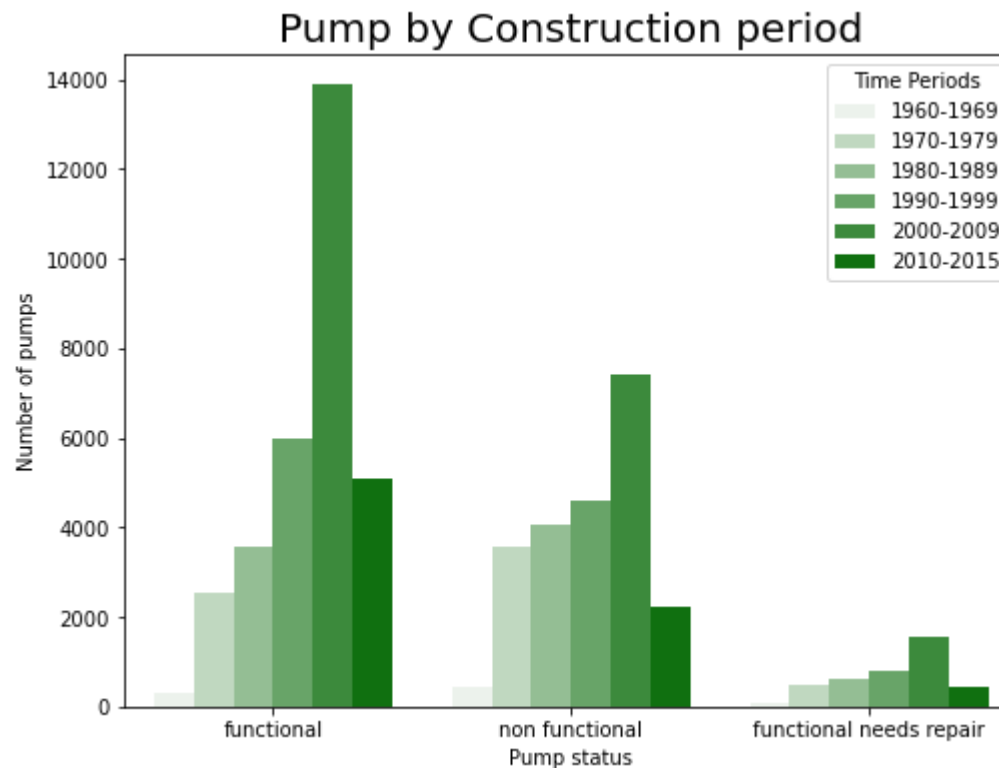
construction_periods	1960-1969	1970-1979	1980-1989	1990-1999	2000-2009	2010-2015
status_group						
functional	283	2542	3579	6002	13893	5090
functional needs repair	69	503	600	784	1553	422
non functional	447	3558	4044	4598	7403	2218

```
In [50]: ▶ # sns.set_style('darkgrid')
plt.figure(figsize=(8,6))
sns.countplot(data=df2, x='status_group', hue='construction_periods', color='Green')

plt.title('Pump by Construction period', fontsize=20)
plt.xlabel('Pump status')
plt.ylabel('Number of pumps')
plt.xticks(rotation = 0)
plt.legend(title = 'Time Periods')
```

executed in 311ms, finished 08:10:20 2021-07-11

Out[50]: <matplotlib.legend.Legend at 0x1ef8b0429a0>



```
In [51]: df2.population.value_counts(normalize=True)
```

executed in 14ms, finished 08:10:20 2021-07-11

```
Out[51]: 0      0.339810
         1      0.121987
         200    0.033688
         150    0.032854
         250    0.029190
         ...
         3241   0.000017
         1960   0.000017
         1685   0.000017
         2248   0.000017
         1439   0.000017
         Name: population, Length: 1049, dtype: float64
```

The population column has more than 33% of its data to be 0 values and 12% of only 1 value. This information doesn't seem correct at all, better drop that feature. Some of the features are repetitive so we will just use the one with the most data and drop all the other similar features.


```
In [52]: df2.drop(columns=['wpt_name', 'recorded_by', 'lga', 'ward',
                           'extraction_type', 'extraction_type_group', 'population',
                           'management', 'payment_type', 'water_quality', 'source',
                           'source_class', 'waterpoint_type', 'quantity_group',
                           'region_code', 'const_year', 'district_code'], axis=1, inplace=True)
```

executed in 15ms, finished 08:10:20 2021-07-11

In [53]:  df2.shape

executed in 15ms, finished 08:10:20 2021-07-11

Out[53]: (57588, 19)

In [54]:  *# Replace target values - there are three classes*
df2 = df2.replace({'status_group': {'functional' : 1,
 'non functional' : 0,
 'functional needs repair' : 2}})
Check to see that it worked
df2.iloc[15:20]

executed in 62ms, finished 08:10:21 2021-07-11

Out[54]:

	id	amount_tsh	funder	gps_height	installer	longitude	latitude	basin	region	scheme_management	extraction_ty
15	61848	0.0	Rudep	1645	DWE	31.444121	-8.274962	Lake Tanganyika	Rukwa	VWC	h
16	48451	500.0	Unicef	1703	DWE	34.642439	-9.106185	Rufiji	Iringa	WUA	
17	58155	0.0	Unicef	1656	DWE	34.569266	-9.085515	Rufiji	Iringa	WUA	
18	34169	0.0	Hesawa	1162	DWE	32.920154	-1.947868	Lake Victoria	Mwanza	other	
19	18274	500.0	Danida	1763	Danid	34.508967	-9.894412	Lake Nyasa	Iringa	VWC	

1.5.4 Creating Dummies

```
In [55]: target = ['status_group']

categorical = ['funder', 'installer', 'basin', 'region',
               'scheme_management', 'extraction_type_class', 'management_group', 'payment',
               'quality_group', 'quantity', 'source_type', 'waterpoint_type_group',
               'construction_periods']

continuous = ['amount_tsh', 'gps_height', 'longitude', 'latitude']

executed in 14ms, finished 08:10:21 2021-07-11
```

```
In [56]: # print number of classes in each category
for col in categorical:
    print(col, df2[col].value_counts().count())

executed in 92ms, finished 08:10:21 2021-07-11
```

```
funder 1859
installer 2114
basin 9
region 21
scheme_management 13
extraction_type_class 7
management_group 5
payment 7
quality_group 6
quantity 5
source_type 7
waterpoint_type_group 6
construction_periods 6
```

```
In [57]: categories_to_remove = {}
for col in categorical:
    df_tmp = pd.DataFrame(df2[col].value_counts(normalize=True))
    other_categories = list(df_tmp.loc[df_tmp[col]<0.01].index)
    df2[col] = df2[col].map(lambda x: 'other' if x in other_categories else x)
    categories_to_remove[col] = other_categories

executed in 1.46s, finished 08:10:22 2021-07-11
```

```
In [58]: # checking if our features of less than 1% are replaced
for col in categorical:
    print(df2[col].value_counts(normalize=True), '\n', df2[col].value_counts(normalize=True).count(), '\n')
```

executed in 155ms, finished 08:10:22 2021-07-11

other	0.527593
Government Of Tanzania	0.153539
Danida	0.054074
Hesawa	0.033236
World Bank	0.023356
Kkkt	0.022348
World Vision	0.021254
Rwssp	0.020612
Unicef	0.017972
District Council	0.014638
Tasaf	0.014482
Dhv	0.014395
Private Individual	0.014309
0	0.013492
Norad	0.013284
Germany Republi	0.010592
Tcrs	0.010454
Ministry Of Water	0.010245
Water	0.010124

```
In [59]: df2.drop(target , axis=1).columns
```

executed in 15ms, finished 08:10:22 2021-07-11

```
Out[59]: Index(['id', 'amount_tsh', 'funder', 'gps_height', 'installer', 'longitude',
               'latitude', 'basin', 'region', 'scheme_management',
               'extraction_type_class', 'management_group', 'payment', 'quality_group',
               'quantity', 'source_type', 'waterpoint_type_group',
               'construction_periods'],
              dtype='object')
```


```
In [60]: df_dummies = pd.get_dummies(df2.drop(target, axis=1))
df_dummies.head().sort_values(by='id')
```

executed in 111ms, finished 08:10:22 2021-07-11

Out[60]:

	id	amount_tsh	gps_height	longitude	latitude	funder_0	funder_Danida	funder_Dhv	funder_District Council	funder_Germany Republi	...
1	8776	0.0	1399	34.698766	-2.147466	0	0	0	0	0	...
4	19728	0.0	0	31.130847	-1.825359	0	0	0	0	0	...
2	34310	25.0	686	37.460664	-3.821329	0	0	0	0	0	...
3	67743	0.0	263	38.486161	-11.155298	0	0	0	0	0	...
0	69572	6000.0	1390	34.938093	-9.856322	0	0	0	0	0	...

5 rows × 119 columns



```
In [61]: df_dummies.shape
```

executed in 15ms, finished 08:10:22 2021-07-11

Out[61]: (57588, 119)

The next step would be to concatenate the Target features with the dummies.

```
In [62]: df_dummies['status_group'] = df2['status_group'].values
```

executed in 14ms, finished 08:10:22 2021-07-11

```
In [63]: df_dummies.shape
```

executed in 15ms, finished 08:10:22 2021-07-11

Out[63]: (57588, 120)

1.6 Modeling

Here, we will run some Models using the classification algorithms of KNN, Random Forest and XGBoost. First we will run baseline models in each method and then move on to tuning and optimizing those models to increase performance and metric scores. I will use F1-score as my deciding metric, but precision and recall will let us know what values we're having trouble classifying, and where I can improve.

```
In [64]: ▶ # assign variables for features and target
X = df_dummies.drop('status_group', axis = 1)
y = df_dummies['status_group']
```

executed in 30ms, finished 08:10:22 2021-07-11

```
In [65]: ▶ # split into test and train data sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state = 42)
```

executed in 47ms, finished 08:10:22 2021-07-11

```
In [66]: ▶ len(y_test)
```

executed in 14ms, finished 08:10:23 2021-07-11

Out[66]: 11518

```
In [67]: ▶ from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score

models = []
def get_metrics(y_test, X_test, model):
    labels = y_test.to_numpy()
    preds = model.predict(X_test)

    metrics = {}
    metrics['accuracy'] = accuracy_score(labels, preds)
    metrics['f1'] = f1_score(labels, preds, average='macro')
    metrics['precision'] = precision_score(labels, preds, average='macro')
    metrics['recall'] = recall_score(labels, preds, average='macro')

    return metrics
```

executed in 13ms, finished 08:10:23 2021-07-11

1.7 KNN

The KNN model is simple to fit, but time-consuming to predict on, especially on this large dataset. It also has relatively few hyperparameters to tune, so it may not improve much.

1.7.1 Baseline Model

In [68]: `from sklearn.neighbors import KNeighborsClassifier`

```
knn = KNeighborsClassifier()
knn.fit(X_train, y_train);

metrics = get_metrics(y_test, X_test, knn)
metrics['name'] = 'K Nearest Neighbors Baseline'
models.append(metrics)
```

executed in 4.05s, finished 08:10:27 2021-07-11

In [69]: `print(classification_report(y_train, knn.predict(X_train),
target_names=['nonfunctional', 'functional',
'needs repair']))`

executed in 5.17s, finished 08:10:32 2021-07-11

	precision	recall	f1-score	support
nonfunctional	0.61	0.65	0.63	17842
functional	0.71	0.77	0.74	25088
needs repair	0.61	0.04	0.07	3140
accuracy			0.67	46070
macro avg	0.65	0.48	0.48	46070
weighted avg	0.67	0.67	0.65	46070

```
In [70]: ▶ print(confusion_matrix(y_test, knn.predict(X_test)))
print(classification_report(y_test, knn.predict(X_test)))
```

executed in 2.72s, finished 08:10:34 2021-07-11

```
[[2092 2312   22]
 [2433 3843   25]
 [ 317  468    6]]
      precision    recall  f1-score   support

     0       0.43      0.47      0.45       4426
     1       0.58      0.61      0.59       6301
     2       0.11      0.01      0.01        791

 accuracy          0.52       11518
 macro avg          0.38      0.36      0.35       11518
 weighted avg          0.49      0.52      0.50       11518
```

1.7.2 Confusion matrix


```
In [71]: ▶ plt.figure(figsize=(6,5))
plot_confusion_matrix(knn, X_test, y_test,
                      display_labels=['nonfunctional', 'functional',
                                     'needs repair'], values_format='d',
                      cmap=plt.cm.Blues)

plt.title('K Nearest Neighbors')
plt.tight_layout()

plt.show()
```

executed in 1.56s, finished 08:10:36 2021-07-11

<Figure size 432x360 with 0 Axes>



In this KNN our model was able to capture 47% of the non functional, 61% of the functional, and only 1% of the functional needs repair.

1.7.3 Standardizing

First scale the data. We scale the data after splitting the train and test data to avoid data leakage.


```

In [72]: # Import StandardScaler
from sklearn.preprocessing import StandardScaler

# Instantiate StandardScaler
scaler = StandardScaler()

# Transform the training and test sets
scaled_data_train = scaler.fit_transform(X_train)
scaled_data_test = scaler.transform(X_test)

# Convert into a DataFrame
scaled_df_train = pd.DataFrame(scaled_data_train, columns=X.columns)
scaled_df_train.head()


```

executed in 173ms, finished 08:10:36 2021-07-11

Out[72]:

	id	amount_tsh	gps_height	longitude	latitude	funder_0	funder_Danida	funder_Dhv	funder_District Council	funder_Germany Republi
0	1.273710	2.393683	1.172595	0.057318	-1.376628	-0.11822	-0.240045	-0.120373	-0.122854	-0.102501
1	-1.251616	-0.104170	1.035817	-0.205919	0.330106	-0.11822	-0.240045	-0.120373	-0.122854	-0.102501
2	-0.368808	-0.104170	0.474307	0.603923	0.441211	-0.11822	-0.240045	-0.120373	-0.122854	-0.102501
3	-1.426384	-0.104170	-0.994257	-1.321809	1.661788	-0.11822	-0.240045	-0.120373	-0.122854	-0.102501
4	0.708950	-0.104170	0.838569	-1.917170	0.502559	-0.11822	-0.240045	-0.120373	-0.122854	-0.102501

5 rows × 119 columns



Now that you've preprocessed the data it's time to train a KNN classifier and validate its accuracy.

```
In [73]: ▶ # Import KNeighborsClassifier
from sklearn.neighbors import KNeighborsClassifier

# Instantiate KNeighborsClassifier
knn = KNeighborsClassifier()
# Fit the classifier
knn.fit(scaled_data_train, y_train);

# Predict on the test set

metrics = get_metrics(y_test, scaled_data_test, knn)
metrics['name'] = 'K Nearest Neighbors scaled'
models.append(metrics)
```

executed in 1m 2.34s, finished 08:11:39 2021-07-11

1.7.4 Optimazing k-value

Let's first create a function to iterate over a range of K-values to find out the best value for the optimum f1-score. Then pass that value for the first round of the GridSearchCV and take note of the result. In the second round of the GridSearchCV, we will try and narrow down the values around the successful ones already found in the first pass.

1.7.5 Grid search CV

```
In [74]: ▶ def find_best_k(X_train, y_train, X_test, y_test, min_k=1, max_k=25):
    best_k = 0
    best_score = 0.0
    for k in range(min_k, max_k+1, 2):
        knn = KNeighborsClassifier(n_neighbors=k)
        knn.fit(X_train, y_train)
        preds = knn.predict(X_test)
        f1 = f1_score(y_test, preds, average='weighted')
        if f1 > best_score:
            best_k = k
            best_score = f1

    print("Best Value for k: {}".format(best_k))
    print("F1-Score: {}".format(best_score))
```

executed in 13ms, finished 08:11:39 2021-07-11

```
In [75]: ▶ find_best_k(scaled_data_train, y_train, scaled_data_test, y_test)
```

executed in 14m 9s, finished 08:25:48 2021-07-11

Best Value for k: 7
F1-Score: 0.7433163911145209

```
In [76]: ▶ from sklearn.model_selection import GridSearchCV
```

executed in 14ms, finished 08:25:48 2021-07-11

```
In [77]:  param_grid = {
            'n_neighbors': [1, 5, 9], # default 5
            'weights': ['uniform', 'distance'], # default 'uniform'
            'leaf_size': [30, 40], # default 30
            'p': [1, 2] # default 2
        }

        grid_search = GridSearchCV(knn, param_grid, cv=3, scoring='f1_macro')
        grid_search.fit(scaled_data_train, y_train)
```

executed in 1h 1m 45s, finished 09:27:33 2021-07-11

```
Out[77]:  GridSearchCV(cv=3, error_score=nan,
                      estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                                    metric='minkowski',
                                                    metric_params=None, n_jobs=None,
                                                    n_neighbors=5, p=2,
                                                    weights='uniform'),
                      iid='deprecated', n_jobs=None,
                      param_grid={'leaf_size': [30, 40], 'n_neighbors': [1, 5, 9],
                                   'p': [1, 2], 'weights': ['uniform', 'distance']},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring='f1_macro', verbose=0)
```

```
In [78]:  grid_search.best_params_
```

executed in 14ms, finished 09:27:33 2021-07-11

```
Out[78]:  {'leaf_size': 30, 'n_neighbors': 9, 'p': 1, 'weights': 'distance'}
```

```
In [79]:  knn_tuned = KNeighborsClassifier(n_neighbors=9, weights='distance', leaf_size=30, p=1)
        knn_tuned.fit(scaled_data_train, y_train)
        get_metrics(y_test, scaled_data_test, knn_tuned)
        metrics['name'] = 'K Nearest Neighbors tuned'
        models.append(metrics)
```

executed in 1m 2.70s, finished 09:28:36 2021-07-11

```
In [80]: ▶ print(classification_report(y_train, knn_tuned.predict(scaled_data_train),
                                     target_names=['nonfunctional', 'functional',
                                     'needs repair']))
```

executed in 3m 55s, finished 09:32:30 2021-07-11

	precision	recall	f1-score	support
nonfunctional	1.00	1.00	1.00	17842
functional	1.00	1.00	1.00	25088
needs repair	1.00	1.00	1.00	3140
accuracy			1.00	46070
macro avg	1.00	1.00	1.00	46070
weighted avg	1.00	1.00	1.00	46070

```
In [81]: ▶ print(confusion_matrix(y_test, knn_tuned.predict(scaled_data_test)))
print(classification_report(y_test, knn_tuned.predict(scaled_data_test),
                                     target_names=['nonfunctional', 'functional',
                                     'needs repair']))
```

executed in 1m 58.9s, finished 09:34:29 2021-07-11

```
[[3149 1175 102]
 [ 764 5339 198]
 [ 160  427 204]]
```

	precision	recall	f1-score	support
nonfunctional	0.77	0.71	0.74	4426
functional	0.77	0.85	0.81	6301
needs repair	0.40	0.26	0.32	791
accuracy			0.75	11518
macro avg	0.65	0.61	0.62	11518
weighted avg	0.75	0.75	0.75	11518

```
In [82]: # Second round of GridSearchCV selection of parameters
param_grid = {
    'n_neighbors': [8, 9, 10], # default 5
    # 'leaf_size': [30] # default 30
}
knn = KNeighborsClassifier()
grid_search = GridSearchCV(knn, param_grid, cv=3, scoring='f1_macro')
grid_search.fit(scaled_data_train, y_train)
```

executed in 8m 47s, finished 09:43:16 2021-07-11

```
Out[82]: GridSearchCV(cv=3, error_score=nan,
                      estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                                    metric='minkowski',
                                                    metric_params=None, n_jobs=None,
                                                    n_neighbors=5, p=2,
                                                    weights='uniform'),
                      iid='deprecated', n_jobs=None,
                      param_grid={'n_neighbors': [8, 9, 10]}, pre_dispatch='2*n_jobs',
                      refit=True, return_train_score=False, scoring='f1_macro',
                      verbose=0)
```

```
In [83]: grid_search.best_params_
```

executed in 13ms, finished 09:43:16 2021-07-11

```
Out[83]: {'n_neighbors': 9}
```

Since we don't see any change in the parameters selection we are going to keep our first trial and proceed to the next step.

```
In [84]: ▶ print(confusion_matrix(y_test, knn_tuned.predict(scaled_data_test)))
print(classification_report(y_test, knn_tuned.predict(scaled_data_test),
                        target_names=['nonfunctional', 'functional',
                                      'needs repair']))
```

executed in 1m 56.9s, finished 09:45:13 2021-07-11

```
[[3149 1175 102]
 [ 764 5339 198]
 [ 160  427 204]]

              precision    recall  f1-score   support

nonfunctional      0.77        0.71        0.74        4426
functional         0.77        0.85        0.81        6301
needs repair       0.40        0.26        0.32         791

accuracy                   0.75        11518
macro avg                 0.65        0.61        0.62        11518
weighted avg              0.75        0.75        0.75        11518
```

1.7.6 Confusion Matrix

```
In [85]: ▶ from sklearn.metrics import plot_confusion_matrix
```


executed in 13ms, finished 09:45:13 2021-07-11

```
In [86]: plot_confusion_matrix(knn_tuned, X_test, y_test,
                                display_labels=['nonfunctional', 'functional', 'needs repair'],
                                values_format='d', cmap=plt.cm.Blues)

plt.title('K Nearest Neighbors tuned')
plt.tight_layout()

plt.show()
```

executed in 1m 29.5s, finished 09:46:43 2021-07-11



From the confusion matrix we can conclude that using the Grid search CV tuning method our model was able to predict:

Non functional

- 120 True positives
- $(45+6) = 51$ False positives
- $(2300+75) = 2375$ False negatives
- $(3900+33+460+8) = 6947$ True Negatives

Functional

- 3900 True positives
- $(2300+460) = 2760$ False positives
- $(45+33) = 78$ False negatives
- $(120+75+6+8) = 209$ True negatives

Needs Repair

- 8 True positives
- $(75+33) = 108$ False positive
- $(460+6) = 466$ False negative
- $(120+2300+3900+45) = 6365$ True Negative

Since we are more interested in less False positives especially for identifying the 'nonfunctional' and 'needs repair' wells, our KNN model seems to have done better than the baseline model. It still looks like it needs some more work done at correctly classifying the 'needs repair' class, but this could also be due to the fact that the classes 'nonfunctional' and 'needs repair' have more or less similar features affecting their functionality.

We'll try another optimizing technique and see how well our KNN model would perform.

1.7.7 Smote

```
In [87]:  from imblearn.over_sampling import SMOTE
```

```
executed in 203ms, finished 09:46:43 2021-07-11
```

```
In [88]: ▶ # Previous original class distribution
print('Original class distribution: \n')
print(y.value_counts())
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_sample(scaled_data_train, y_train)
# Preview synthetic sample class distribution
print('-----')
print('Synthetic sample class distribution: \n')
print(pd.Series(y_train_resampled).value_counts())
```

executed in 41.0s, finished 09:47:24 2021-07-11

Original class distribution:

```
1    31389
0    22268
2     3931
```

Name: status_group, dtype: int64

Synthetic sample class distribution:

```
2    25088
1    25088
0    25088
```

Name: status_group, dtype: int64

```
In [89]: ▶ knn_smote = KNeighborsClassifier(n_neighbors=9, weights='distance',
                                           leaf_size=30, p=1)
knn_smote.fit(X_train_resampled, y_train_resampled)
get_metrics(y_test, scaled_data_test, knn_smote)
metrics['name'] = 'K Nearest Neighbors smote'
models.append(metrics)
```

executed in 1m 36.7s, finished 09:49:01 2021-07-11

```
In [90]: ▶ print(classification_report(y_train_resampled, knn_smote.predict(X_train_resampled),
                                     target_names=['nonfunctional', 'functional',
                                     'needs repair']))
```

executed in 9m 7s, finished 09:58:07 2021-07-11

	precision	recall	f1-score	support
nonfunctional	1.00	1.00	1.00	25088
functional	1.00	1.00	1.00	25088
needs repair	1.00	1.00	1.00	25088
accuracy			1.00	75264
macro avg	1.00	1.00	1.00	75264
weighted avg	1.00	1.00	1.00	75264

```
In [91]: ▶ print(confusion_matrix(y_test, knn_smote.predict(scaled_data_test)))
print(classification_report(y_test, knn_smote.predict(scaled_data_test),
                                     target_names=['nonfunctional', 'functional',
                                     'needs repair']))
```

executed in 2m 57s, finished 10:01:04 2021-07-11

```
[[3162  942  322]
 [ 828 4782  691]
 [ 136  272  383]]
```

	precision	recall	f1-score	support
nonfunctional	0.77	0.71	0.74	4426
functional	0.80	0.76	0.78	6301
needs repair	0.27	0.48	0.35	791
accuracy			0.72	11518
macro avg	0.61	0.65	0.62	11518
weighted avg	0.75	0.72	0.73	11518

In []: ▶


1.7.8 Confusion Matrix

```
In [92]: plot_confusion_matrix(knn_smote, scaled_data_test, y_test,
                                display_labels=['nonfunctional', 'functional',
                                                'needs repair'],
                                values_format='d', cmap=plt.cm.Blues)

plt.title('K Nearest Neighbors smote')
plt.tight_layout()

plt.show()
```

executed in 1m 29.7s, finished 10:02:34 2021-07-11



The K Nearest Neighbors model performed better after hyperparameter tuning in regards to Accuracy(74%) and F1 score, but didn't improve significantly with smote. The F1 score of (75%, 80%, 38%) show that there is a class imbalance and hence its effect is visible. But generally KNN also takes in incredibly long time to run on this large dataset, making it impractical to tune further. The model was generally able to capture the following percentiles of the actual wells.

K Nearest Neighbor:

1. Baseline model

Non functional (47%)
Functional (61%)
Functional needs repair (1%)

2. Gridsearch CV

Non functional (71%)
Functional (84%)
Functional needs repair (27%)

3. SMOTE

Non functional (71%)
Functional (76%)
Functional needs repair (41%)

1.8 Random forest

1.8.1 Baseline Model

```
In [93]: ▶ from sklearn.ensemble import RandomForestClassifier

forest = RandomForestClassifier(random_state=12)
forest.fit(X_train, y_train);

metrics = get_metrics(y_test, X_test, forest)
metrics['name'] = 'Random Forest Baseline'
models.append(metrics)
```

executed in 7.50s, finished 10:02:41 2021-07-11

```
In [94]: ▶ print(classification_report(y_train, forest.predict(X_train),
                                     target_names=['nonfunctional', 'functional',
                                     'needs repair']))
```

executed in 1.09s, finished 10:02:42 2021-07-11

	precision	recall	f1-score	support
nonfunctional	1.00	1.00	1.00	17842
functional	1.00	1.00	1.00	25088
needs repair	1.00	1.00	1.00	3140
accuracy			1.00	46070
macro avg	1.00	1.00	1.00	46070
weighted avg	1.00	1.00	1.00	46070

```
In [95]: ▶ print(confusion_matrix(y_test, forest.predict(X_test)))
print(classification_report(y_test, forest.predict(X_test),
                           target_names=['nonfunctional', 'functional',
                           'needs repair']))
```

executed in 602ms, finished 10:02:43 2021-07-11


```
[[3391  959   76]
 [ 598 5543  160]
 [ 147  423  221]]
```

	precision	recall	f1-score	support
nonfunctional	0.82	0.77	0.79	4426
functional	0.80	0.88	0.84	6301
needs repair	0.48	0.28	0.35	791
accuracy			0.79	11518
macro avg	0.70	0.64	0.66	11518
weighted avg	0.79	0.79	0.79	11518

1.8.2 Confusion matrix

```
In [96]: plot_confusion_matrix(forest, X_test, y_test,  
                                display_labels=['nonfunctional', 'functional',  
                                                'needs repair'], values_format='d',  
                                cmap=plt.cm.Blues)  
  
plt.title('Random forest Baseline')  
plt.tight_layout()  
  
plt.show()
```

executed in 506ms, finished 10:02:44 2021-07-11



1.8.3 Standardized

```
In [97]: ▶ # Instantiate RandomForestClassifier  
forest = RandomForestClassifier()  
# Fit the classifier  
forest.fit(scaled_data_train, y_train);  
  
# Predict on the test set  
metrics = get_metrics(y_test, scaled_data_test, forest)  
metrics['name'] = 'Random forest scaled'  
models.append(metrics)
```

executed in 7.39s, finished 10:02:51 2021-07-11

1.8.4 GridSearch CV


```
In [98]: param_grid = {
    'n_estimators': [100, 200], # default 100 #boosting stages
    'max_depth': [30, 35, 40], # default None
    'max_features': [50, 60], # default 'auto': auto=sqrt(# of features)=11, None=# of features=122
    # 'min_samples_split' : [20,30,40],
    # 'min_samples_leaf' : [5, 10, 15]
    # default 'auto': auto=sqrt(# of features)=11, None=# of features=122
}
forest = RandomForestClassifier()
grid_search = GridSearchCV(forest, param_grid, cv=3, scoring='f1_macro')
grid_search.fit(scaled_data_train, y_train)
```

executed in 13m 24s, finished 10:16:15 2021-07-11

```
Out[98]: GridSearchCV(cv=3, error_score=nan,
                    estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                    class_weight=None,
                                                    criterion='gini', max_depth=None,
                                                    max_features='auto',
                                                    max_leaf_nodes=None,
                                                    max_samples=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    n_estimators=100, n_jobs=None,
                                                    oob_score=False,
                                                    random_state=None, verbose=0,
                                                    warm_start=False),
                    iid='deprecated', n_jobs=None,
                    param_grid={'max_depth': [30, 35, 40], 'max_features': [50, 60],
                                'n_estimators': [100, 200]},
                    scoring='f1_macro', verbose=0)
```

```
In [99]: grid_search.best_params_
```

executed in 14ms, finished 10:16:15 2021-07-11

```
Out[99]: {'max_depth': 35, 'max_features': 50, 'n_estimators': 200}
```

```
In [100]: ► forest_tuned = RandomForestClassifier(n_estimators=200, max_depth=35,
max_features=60, min_samples_leaf=35,
min_samples_split=70)

forest_tuned.fit(scaled_data_train, y_train)
get_metrics(y_test, scaled_data_test, forest_tuned)
metrics['name'] = 'Random Forest tuned1'
models.append(metrics)
```

executed in 40.2s, finished 10:16:55 2021-07-11

```
In [101]: ► print(confusion_matrix(y_test, forest_tuned.predict(scaled_data_test)))
print(classification_report(y_test, forest_tuned.predict(scaled_data_test),
target_names=['nonfunctional', 'functional',
'needs repair']))
```

executed in 653ms, finished 10:16:56 2021-07-11

```
[[3042 1370  14]
 [ 499 5778  24]
 [ 145  561  85]]
```

	precision	recall	f1-score	support
nonfunctional	0.83	0.69	0.75	4426
functional	0.75	0.92	0.82	6301
needs repair	0.69	0.11	0.19	791
accuracy			0.77	11518
macro avg	0.76	0.57	0.59	11518
weighted avg	0.77	0.77	0.75	11518

```
In [102]: ▶ param_grid = {
    'n_estimators': [150, 200], # default 100
    'max_depth': [35, 45, 50], # default None
    'max_features': [55, 60, 65],
    # 'min_samples_split' : 70,
    # 'min_samples_leaf' : 35
    # we assume 5 would be the min sample leaf and avoid further search
}
forest = RandomForestClassifier()
grid_search = GridSearchCV(forest, param_grid, cv=3, scoring='f1_macro')
grid_search.fit(scaled_data_train, y_train)
```

executed in 24m 39s, finished 10:41:35 2021-07-11

```
Out[102]: GridSearchCV(cv=3, error_score=nan,
    estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
    class_weight=None,
    criterion='gini', max_depth=None,
    max_features='auto',
    max_leaf_nodes=None,
    max_samples=None,
    min_impurity_decrease=0.0,
    min_impurity_split=None,
    min_samples_leaf=1,
    min_samples_split=2,
    min_weight_fraction_leaf=0.0,
    n_estimators=100, n_jobs=None,
    oob_score=False,
    random_state=None, verbose=0,
    warm_start=False),
    iid='deprecated', n_jobs=None,
    param_grid={'max_depth': [35, 45, 50],
    'max_features': [55, 60, 65],
    'n_estimators': [150, 200]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring='f1_macro', verbose=0)
```

```
In [103]: ▶ grid_search.best_params_
```

executed in 14ms, finished 10:41:35 2021-07-11

```
Out[103]: {'max_depth': 45, 'max_features': 65, 'n_estimators': 150}
```

```
In [104]: ► forest_tuned = RandomForestClassifier(n_estimators=200, max_depth=50,
max_features=55, min_samples_leaf=35,
min_samples_split=70)
forest_tuned.fit(scaled_data_train, y_train)
get_metrics(y_test, scaled_data_test, forest_tuned)
metrics['name'] = 'Random Forest tuned2'
models.append(metrics)
```

executed in 37.1s, finished 10:42:12 2021-07-11

```
In [105]: ► print(confusion_matrix(y_test, forest_tuned.predict(scaled_data_test)))
print(classification_report(y_test, forest_tuned.predict(scaled_data_test),
target_names=['nonfunctional', 'functional',
'needs repair']))
```

executed in 665ms, finished 10:42:13 2021-07-11

```
[[3040 1370  16]
 [ 491 5784  26]
 [ 149  556  86]]

              precision    recall  f1-score   support


nonfunctional      0.83      0.69      0.75      4426
functional         0.75      0.92      0.83      6301
needs repair       0.67      0.11      0.19       791

   accuracy              0.77      11518
  macro avg              0.75      0.57      0.59      11518
weighted avg              0.77      0.77      0.75      11518
```

1.8.5 Confusion matrix

```
In [106]: plot_confusion_matrix(forest_tuned, scaled_data_test, y_test,  
                                display_labels=['nonfunctional', 'functional',  
                                                'needs repair'], values_format='d',  
                                cmap=plt.cm.Blues)  
plt.title('Random Forest Scaled')  
plt.tight_layout()  
  
plt.show()
```

executed in 517ms, finished 10:42:13 2021-07-11



1.8.6 Smote

```
In [107]: ► forest_smote = RandomForestClassifier(n_estimators=150, max_depth=40,
max_features=50, min_samples_leaf=35,
min_samples_split=70)
forest_smote.fit(X_train_resampled, y_train_resampled)
get_metrics(y_test, scaled_data_test, forest_smote)
metrics['name'] = 'Random Forest smote'
models.append(metrics)
```

executed in 47.3s, finished 10:43:01 2021-07-11

```
In [108]: ► print(confusion_matrix(y_test, forest_smote.predict(scaled_data_test)))
print(classification_report(y_test, forest_smote.predict(scaled_data_test),
target_names=['nonfunctional', 'functional',
'needs repair']))
```

executed in 596ms, finished 10:43:01 2021-07-11

```
[[3082  988  356]
 [ 544 5008  749]
 [  86  260  445]]

              precision    recall  f1-score   support


nonfunctional      0.83      0.70      0.76      4426
functional         0.80      0.79      0.80      6301
needs repair       0.29      0.56      0.38       791

   accuracy              0.74      11518
  macro avg              0.64      0.68      0.65      11518
weighted avg              0.78      0.74      0.75      11518
```

1.8.7 Confusion matrix

```
In [109]: plot_confusion_matrix(forest_tuned, scaled_data_test, y_test,  
                                display_labels=['nonfunctional', 'functional',  
                                                'needs repair'], values_format='d',  
                                cmap=plt.cm.Blues)  
plt.title('Random Forest Scaled')  
plt.tight_layout()  
  
plt.show()
```

executed in 579ms, finished 10:43:02 2021-07-11



Our Random Forest model has been able to correctly classify 5,353 data sets out of 6,837 which is 78% of our total data.


2 Feature importance

```
In [110]: # Plot the feature importance of each feature
features = X_train.shape[1]
plt.figure(figsize=(20,20))
plt.barh(range(features), forest_tuned.feature_importances_, align='center')
plt.yticks(np.arange(features), X_train.columns.values)
plt.title('Feature Importance', fontsize=20)
plt.xlabel('Feature importance')
plt.ylabel('Features')
```

executed in 2.52s, finished 10:43:05 2021-07-11

Out[110]: Text(0, 0.5, 'Features')

Feature Importance



2.1 XGBoost

2.1.1 Baseline Model

In []: ▶

In [111]: ▶ `from sklearn.ensemble import GradientBoostingClassifier`

```
# Instantiate XGBClassifier and fit classifier
xgb = XGBClassifier(random_state=12)
xgb.fit(X_train, y_train);

metrics = get_metrics(y_test, X_test, xgb)
metrics['name'] = 'XG Boost Baseline'
models.append(metrics)
```

executed in 11.6s, finished 10:43:16 2021-07-11

C:\Users\milen\anaconda3\lib\site-packages\xgboost\sklearn.py:888: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

warnings.warn(label_encoder_deprecation_msg, UserWarning)

[10:43:05] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
In [112]: ▶ print(classification_report(y_train, xgb.predict(X_train),
                                         target_names=['nonfunctional', 'functional',
                                                         'needs repair']))
```

executed in 202ms, finished 10:43:16 2021-07-11

	precision	recall	f1-score	support
nonfunctional	0.89	0.77	0.82	17842
functional	0.80	0.95	0.87	25088
needs repair	0.86	0.31	0.45	3140
accuracy			0.83	46070
macro avg	0.85	0.67	0.72	46070
weighted avg	0.84	0.83	0.82	46070

```
In [113]: ▶ print(confusion_matrix(y_test, xgb.predict(X_test)))
print(classification_report(y_test, xgb.predict(X_test),
                             target_names=['nonfunctional', 'functional',
                                             'needs repair']))
```

executed in 138ms, finished 10:43:17 2021-07-11


```
[[3157 1237  32]
 [ 491 5748  62]
 [ 126  502 163]]
```

	precision	recall	f1-score	support
nonfunctional	0.84	0.71	0.77	4426
functional	0.77	0.91	0.83	6301
needs repair	0.63	0.21	0.31	791
accuracy			0.79	11518
macro avg	0.75	0.61	0.64	11518
weighted avg	0.78	0.79	0.77	11518

2.1.2 Confusion Matrix

```
In [114]: plot_confusion_matrix(xgb, X_test, y_test,  
                                display_labels=['nonfunctional', 'functional',  
                                                'needs repair'], values_format='d',  
                                cmap=plt.cm.Blues)  
  
plt.title('XGBoost')  
plt.tight_layout()  
  
plt.show()
```

executed in 301ms, finished 10:43:17 2021-07-11



2.1.3 Standardized

```
In [115]: ▶ # Fit the classifier  
xgb.fit(scaled_data_train, y_train);  
  
# Predict on the test set  
metrics = get_metrics(y_test, scaled_data_test, xgb)  
metrics['name'] = 'XGBoost scaled'  
models.append(metrics)
```

executed in 13.6s, finished 10:43:30 2021-07-11

[10:43:17] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

2.1.4 Grid search CV

```
In [116]: ▶ param_grid = {  
    'learning_rate': [0.05, 0.1],  
    'max_depth': [6, 8, 10],  
    'subsample': [0.5, 0.7],  
    'n_estimators': [100, 150, 200],  
}  
xgb = XGBClassifier()  
grid_search = GridSearchCV(xgb, param_grid, cv=None, scoring='f1_macro', n_jobs=1)  
grid_search.fit(scaled_data_train, y_train)
```

executed in 1h 48m 58s, finished 12:32:28 2021-07-11

```
In [117]: ▶ grid_search.best_params_
```

executed in 24ms, finished 12:32:28 2021-07-11


```
Out[117]: {'learning_rate': 0.1, 'max_depth': 10, 'n_estimators': 200, 'subsample': 0.7}
```

In [118]:  *# Instantiate XGBClassifier and fit classifier*

```
xgb_tuned = XGBClassifier(n_estimators=200, learning_rate=0.1, max_depth=10,
                          subsample=0.7)
xgb_tuned.fit(scaled_data_train, y_train)
get_metrics(y_test, scaled_data_test, xgb_tuned)
metrics['name'] = 'XGBoost tuned1'
models.append(metrics)
```

executed in 1m 20.8s, finished 12:33:49 2021-07-11

[12:32:29] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

In [119]:  `print(classification_report(y_train, xgb_tuned.predict(scaled_data_train),
 target_names=['nonfunctional', 'functional',
 'needs repair']))`

executed in 463ms, finished 12:33:50 2021-07-11

	precision	recall	f1-score	support
nonfunctional	0.96	0.89	0.92	17842
functional	0.89	0.98	0.94	25088
needs repair	0.96	0.65	0.77	3140
accuracy			0.92	46070
macro avg	0.94	0.84	0.88	46070
weighted avg	0.92	0.92	0.92	46070

```
In [120]: ▶ print(confusion_matrix(y_test, xgb_tuned.predict(scaled_data_test)))
print(classification_report(y_test, xgb_tuned.predict(scaled_data_test),
                        target_names=['nonfunctional', 'functional',
                                      'needs repair']))
```

executed in 231ms, finished 12:33:50 2021-07-11

```
[[3292 1083   51]
 [ 505 5719   77]
 [ 131  473 187]]
```

	precision	recall	f1-score	support
nonfunctional	0.84	0.74	0.79	4426
functional	0.79	0.91	0.84	6301
needs repair	0.59	0.24	0.34	791
accuracy			0.80	11518
macro avg	0.74	0.63	0.66	11518
weighted avg	0.79	0.80	0.79	11518


```
In [121]: param_grid = {
    'learning_rate': [0.1, 0.2],
    'max_depth': [10, 15, 20],
    'subsample': [0.6, 0.7],
    'n_estimators': [200, 250],
}
xgb = XGBClassifier()
grid_search = GridSearchCV(xgb, param_grid, cv=None, scoring='f1_macro', n_jobs=1)
grid_search.fit(scaled_data_train, y_train)
```

executed in 3h 29m 8s, finished 16:02:57 2021-07-11

d with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[12:54:10] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[12:55:29] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[12:56:50] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[12:58:14] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[12:59:59] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[13:01:41] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
In [122]: grid_search.best_params_
```

executed in 12ms, finished 16:02:57 2021-07-11

```
Out[122]: {'learning_rate': 0.1, 'max_depth': 20, 'n_estimators': 200, 'subsample': 0.7}
```

In [123]: **▶** *# Instantiate XGBClassifier and fit classifier*

```
xgb_tuned = XGBClassifier(n_estimators=100, learning_rate=0.1, max_depth=15,
                          subsample=0.5)
xgb_tuned.fit(scaled_data_train, y_train)
get_metrics(y_test, scaled_data_test, xgb_tuned)
metrics['name'] = 'XGBoost tuned2'
models.append(metrics)
```

executed in 53.4s, finished 16:03:51 2021-07-11

[16:02:58] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

In [124]: **▶**

```
print(confusion_matrix(y_test, xgb_tuned.predict(scaled_data_test)))
print(classification_report(y_test, xgb_tuned.predict(scaled_data_test),
                        target_names=['nonfunctional', 'functional',
                                      'needs repair']))
```

executed in 164ms, finished 16:03:51 2021-07-11

```
[[3336 1049   41]
 [ 527 5693   81]
 [ 132  470 189]]
      precision    recall  f1-score   support


nonfunctional      0.84      0.75      0.79      4426
functional         0.79      0.90      0.84      6301
needs repair       0.61      0.24      0.34       791

   accuracy      0.74
  macro avg      0.74
weighted avg      0.79
```

2.1.5 Confusion matrix

```
In [125]: plot_confusion_matrix(xgb_tuned, scaled_data_test, y_test,  
                                display_labels=['nonfunctional', 'functional',  
                                                'needs repair'], values_format='d',  
                                cmap=plt.cm.Blues)  
  
plt.title('XGBoost tuned')  
plt.tight_layout()  
  
plt.show()
```

executed in 316ms, finished 16:03:51 2021-07-11



2.1.6 Smote

```
In [126]: ▶ # Instantiate XGBClassifier and fit classifier
xgb_smote = XGBClassifier(n_estimators=200, learning_rate=0.1, max_depth=10,
                        subsample=0.7)
xgb_smote.fit(X_train_resampled, y_train_resampled)
get_metrics(y_test, scaled_data_test, forest_smote)
metrics['name'] = 'XGBoost smote'
models.append(metrics)
```

executed in 2m 6s, finished 16:05:57 2021-07-11

[16:03:52] WARNING: ..\src\learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
In [127]: ▶ print(confusion_matrix(y_test, xgb_smote.predict(scaled_data_test)))
print(classification_report(y_test, xgb_smote.predict(scaled_data_test),
                        target_names=['nonfunctional', 'functional',
                                      'needs repair']))
```

executed in 267ms, finished 16:05:57 2021-07-11

```
[[3297  953  176]
 [ 578 5332  391]
 [ 122  336  333]]

              precision    recall  f1-score   support

nonfunctional      0.82      0.74      0.78      4426
functional         0.81      0.85      0.83      6301
needs repair       0.37      0.42      0.39       791


   accuracy      0.78
  macro avg      0.67
weighted avg      0.78
```

2.1.7 Confusion matrix

```
In [128]: ▶ plot_confusion_matrix(xgb_smote, scaled_data_test, y_test,
                                display_labels=['nonfunctional', 'functional',
                                                'needs repair'], values_format='d',
                                cmap=plt.cm.Blues)
plt.title('XGBoost smote')
plt.tight_layout()

plt.show()
```

executed in 378ms, finished 16:05:58 2021-07-11



In []: ▶

2.2 Analysis

```
In [129]: models_df = pd.DataFrame(models)
models_df.sort_values(by='accuracy', ascending=False)
```

executed in 30ms, finished 16:05:58 2021-07-11

Out[129]:

	accuracy	f1	precision	recall	name
5	0.797447	0.664382	0.705514	0.644198	Random Forest smote
6	0.797447	0.664382	0.705514	0.644198	Random Forest smote
7	0.797447	0.664382	0.705514	0.644198	Random Forest smote
8	0.797447	0.664382	0.705514	0.644198	Random Forest smote
4	0.794843	0.661490	0.701299	0.641750	Random Forest Baseline
9	0.787289	0.638279	0.746162	0.610530	XG Boost Baseline
10	0.786074	0.634550	0.740820	0.607812	XGBoost smote
11	0.786074	0.634550	0.740820	0.607812	XGBoost smote
12	0.786074	0.634550	0.740820	0.607812	XGBoost smote
13	0.786074	0.634550	0.740820	0.607812	XGBoost smote
1	0.751780	0.609579	0.663459	0.590658	K Nearest Neighbors smote

2.3 Conclusion & Recommendation

The overall best model was Random Forrest which performed considerably better than other models at predicting which wells were functional but need repair.(Recall score)

Random forest:

1. Baseline model

Non functional (77%)

Functional (87%)

Functional needs repair (29%)

2. Gridsearch CV

Non functional (68%)
Functional (92%)
Functional needs repair (10%)

3. SMOTE

Non functional (69%)
Functional (79%)
Functional needs repair (55%)

Given more time and with some more tuning it may be able to increase its performance.

2.4 Further Study

1. Identify and indicate wells that are no longer functional due to being past their life span.
2. Review which extraction types last longer in supply and quality of water.
3. Identify companies with a record of poor installation, and lack of maintenance.
4. Include date records of maintenance, accurate population size and Construction year.

In []: ▶